

# Logistic Regression The Telecom Churn Case Study

August 12, 2020

## 0.1 Telecom Churn Case Study

With 21 predictor variables we need to predict whether a particular customer will switch to another telecom provider or not. In telecom terminology, this is referred to as churning and not churning, respectively.

### 0.1.1 Step 1: Importing and Merging Data

```
In [1]: # Suppressing Warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # Importing Pandas and NumPy
import pandas as pd, numpy as np
```

```
In [3]: # Importing all datasets
churn_data = pd.read_csv("churn_data.csv")
churn_data.head()
```

```
Out[3]:
```

	customerID	tenure	PhoneService	Contract	PaperlessBilling	\
0	7590-VHVEG	1	No	Month-to-month		Yes
1	5575-GNVDE	34	Yes	One year		No
2	3668-QPYBK	2	Yes	Month-to-month		Yes
3	7795-CFOCW	45	No	One year		No
4	9237-HQITU	2	Yes	Month-to-month		Yes

		PaymentMethod	MonthlyCharges	TotalCharges	Churn
0		Electronic check	29.85	29.85	No
1		Mailed check	56.95	1889.5	No
2		Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)		42.30	1840.75	No
4		Electronic check	70.70	151.65	Yes

```
In [4]: customer_data = pd.read_csv("customer_data.csv")
customer_data.head()
```

```
Out[4]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents
0	7590-VHVEG	Female	0	Yes	No

1	5575-GNVDE	Male	0	No	No
2	3668-QPYBK	Male	0	No	No
3	7795-CFOCW	Male	0	No	No
4	9237-HQITU	Female	0	No	No

```
In [5]: internet_data = pd.read_csv("internet_data.csv")
internet_data.head()
```

```
Out [5]:
```

	customerID	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	7590-VHVEG	No phone service	DSL	No	Yes	
1	5575-GNVDE	No	DSL	Yes	No	
2	3668-QPYBK	No	DSL	Yes	Yes	
3	7795-CFOCW	No phone service	DSL	Yes	No	
4	9237-HQITU	No	Fiber optic	No	No	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies
0	No	No	No	No
1	Yes	No	No	No
2	No	No	No	No
3	Yes	Yes	No	No
4	No	No	No	No

## Combining all data files into one consolidated dataframe

```
In [6]: # Merging on 'customerID'
df_1 = pd.merge(churn_data, customer_data, how='inner', on='customerID')
```

```
In [7]: # Final dataframe with all predictor variables
telecom = pd.merge(df_1, internet_data, how='inner', on='customerID')
```

### 0.1.2 Step 2: Inspecting the Dataframe

```
In [8]: # Let's see the head of our master dataset
telecom.head()
```

```
Out [8]:
```

	customerID	tenure	PhoneService	Contract	PaperlessBilling	\
0	7590-VHVEG	1	No	Month-to-month	Yes	
1	5575-GNVDE	34	Yes	One year	No	
2	3668-QPYBK	2	Yes	Month-to-month	Yes	
3	7795-CFOCW	45	No	One year	No	
4	9237-HQITU	2	Yes	Month-to-month	Yes	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn	gender	\
0	Electronic check	29.85	29.85	No	Female	
1	Mailed check	56.95	1889.5	No	Male	
2	Mailed check	53.85	108.15	Yes	Male	
3	Bank transfer (automatic)	42.30	1840.75	No	Male	
4	Electronic check	70.70	151.65	Yes	Female	

	...	Partner	Dependents	MultipleLines	InternetService	\
0	...	Yes	No	No phone service		DSL
1	...	No	No		No	DSL
2	...	No	No		No	DSL
3	...	No	No	No phone service		DSL
4	...	No	No		No	Fiber optic

	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	\
0	No	Yes	No	No	No	
1	Yes	No	Yes	No	No	
2	Yes	Yes	No	No	No	
3	Yes	No	Yes	Yes	No	
4	No	No	No	No	No	

	StreamingMovies
0	No
1	No
2	No
3	No
4	No

[5 rows x 21 columns]

In [9]: *# Let's check the dimensions of the dataframe*  
telecom.shape

Out[9]: (7043, 21)

In [10]: *# let's look at the statistical aspects of the dataframe*  
telecom.describe()

```
Out[10]:
```

	tenure	MonthlyCharges	SeniorCitizen
count	7043.000000	7043.000000	7043.000000
mean	32.371149	64.761692	0.162147
std	24.559481	30.090047	0.368612
min	0.000000	18.250000	0.000000
25%	9.000000	35.500000	0.000000
50%	29.000000	70.350000	0.000000
75%	55.000000	89.850000	0.000000
max	72.000000	118.750000	1.000000

In [11]: *# Let's see the type of each column*  
telecom.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7043 entries, 0 to 7042
Data columns (total 21 columns):
customerID      7043 non-null object
tenure          7043 non-null int64
```

```

PhoneService      7043 non-null object
Contract          7043 non-null object
PaperlessBilling  7043 non-null object
PaymentMethod     7043 non-null object
MonthlyCharges    7043 non-null float64
TotalCharges      7043 non-null object
Churn             7043 non-null object
gender            7043 non-null object
SeniorCitizen     7043 non-null int64
Partner           7043 non-null object
Dependents        7043 non-null object
MultipleLines     7043 non-null object
InternetService   7043 non-null object
OnlineSecurity    7043 non-null object
OnlineBackup      7043 non-null object
DeviceProtection  7043 non-null object
TechSupport       7043 non-null object
StreamingTV       7043 non-null object
StreamingMovies   7043 non-null object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.2+ MB

```

### 0.1.3 Step 3: Data Preparation

#### Converting some binary variables (Yes/No) to 0/1

In [12]: *# List of variables to map*

```

varlist = ['PhoneService', 'PaperlessBilling', 'Churn', 'Partner', 'Dependents']

# Defining the map function
def binary_map(x):
    return x.map({'Yes': 1, "No": 0})

# Applying the function to the housing list
telecom[varlist] = telecom[varlist].apply(binary_map)

```

In [13]: telecom.head()

```

Out[13]:
  customerID  tenure  PhoneService  Contract  PaperlessBilling  \
0  7590-VHVEG      1              0  Month-to-month            1
1  5575-GNVDE     34              1      One year            0
2  3668-QPYBK      2              1  Month-to-month            1
3  7795-CFOCW     45              0      One year            0
4  9237-HQITU      2              1  Month-to-month            1

      PaymentMethod  MonthlyCharges  TotalCharges  Churn  gender  \
0  Electronic check           29.85          29.85     0  Female

```

1		Mailed check	56.95	1889.5	0	Male
2		Mailed check	53.85	108.15	1	Male
3	Bank transfer (automatic)		42.30	1840.75	0	Male
4		Electronic check	70.70	151.65	1	Female

	...	Partner	Dependents	MultipleLines	InternetService	\
0	...	1	0	No phone service	DSL	
1	...	0	0	No	DSL	
2	...	0	0	No	DSL	
3	...	0	0	No phone service	DSL	
4	...	0	0	No	Fiber optic	

	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	\
0	No	Yes	No	No	No	
1	Yes	No	Yes	No	No	
2	Yes	Yes	No	No	No	
3	Yes	No	Yes	Yes	No	
4	No	No	No	No	No	

	StreamingMovies
0	No
1	No
2	No
3	No
4	No

[5 rows x 21 columns]

**For categorical variables with multiple levels, create dummy features (one-hot encoded)**

```
In [14]: # Creating a dummy variable for some of the categorical variables and dropping the first dummy1 = pd.get_dummies(telecom[['Contract', 'PaymentMethod', 'gender', 'InternetService']])
```

```
# Adding the results to the master dataframe
telecom = pd.concat([telecom, dummy1], axis=1)
```

```
In [15]: telecom.head()
```

Out[15]:	customerID	tenure	PhoneService	Contract	PaperlessBilling	\
0	7590-VHVEG	1	0	Month-to-month	1	
1	5575-GNVDE	34	1	One year	0	
2	3668-QPYBK	2	1	Month-to-month	1	
3	7795-CFOCW	45	0	One year	0	
4	9237-HQITU	2	1	Month-to-month	1	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn	gender	\
0	Electronic check	29.85	29.85	0	Female	
1	Mailed check	56.95	1889.5	0	Male	

2	Mailed check	53.85	108.15	1	Male
3	Bank transfer (automatic)	42.30	1840.75	0	Male
4	Electronic check	70.70	151.65	1	Female

	...	StreamingTV	StreamingMovies	Contract_One year	\
0	...	No	No	0	
1	...	No	No	1	
2	...	No	No	0	
3	...	No	No	1	
4	...	No	No	0	

	Contract_Two year	PaymentMethod_Credit card (automatic)	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	PaymentMethod_Electronic check	PaymentMethod_Mailed check	gender_Male	\
0	1	0	0	
1	0	1	1	
2	0	1	1	
3	0	0	1	
4	1	0	0	

	InternetService_Fiber optic	InternetService_No
0	0	0
1	0	0
2	0	0
3	0	0
4	1	0

[5 rows x 29 columns]

In [16]: # Creating dummy variables for the remaining categorical variables and dropping the l

```

# Creating dummy variables for the variable 'MultipleLines'
ml = pd.get_dummies(telecom['MultipleLines'], prefix='MultipleLines')
# Dropping MultipleLines_No phone service column
ml1 = ml.drop(['MultipleLines_No phone service'], 1)
#Adding the results to the master dataframe
telecom = pd.concat([telecom,ml1], axis=1)

# Creating dummy variables for the variable 'OnlineSecurity'.
os = pd.get_dummies(telecom['OnlineSecurity'], prefix='OnlineSecurity')
os1 = os.drop(['OnlineSecurity_No internet service'], 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,os1], axis=1)

```

```

# Creating dummy variables for the variable 'OnlineBackup'.
ob = pd.get_dummies(telecom['OnlineBackup'], prefix='OnlineBackup')
ob1 = ob.drop(['OnlineBackup_No internet service'], 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,ob1], axis=1)

# Creating dummy variables for the variable 'DeviceProtection'.
dp = pd.get_dummies(telecom['DeviceProtection'], prefix='DeviceProtection')
dp1 = dp.drop(['DeviceProtection_No internet service'], 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,dp1], axis=1)

# Creating dummy variables for the variable 'TechSupport'.
ts = pd.get_dummies(telecom['TechSupport'], prefix='TechSupport')
ts1 = ts.drop(['TechSupport_No internet service'], 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,ts1], axis=1)

# Creating dummy variables for the variable 'StreamingTV'.
st =pd.get_dummies(telecom['StreamingTV'], prefix='StreamingTV')
st1 = st.drop(['StreamingTV_No internet service'], 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,st1], axis=1)

# Creating dummy variables for the variable 'StreamingMovies'.
sm = pd.get_dummies(telecom['StreamingMovies'], prefix='StreamingMovies')
sm1 = sm.drop(['StreamingMovies_No internet service'], 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,sm1], axis=1)

```

```
In [17]: telecom.head()
```

```

Out[17]:   customerID  tenure  PhoneService      Contract  PaperlessBilling  \
0  7590-VHVEG      1          0  Month-to-month              1
1  5575-GNVDE     34          1      One year              0
2  3668-QPYBK      2          1  Month-to-month              1
3  7795-CFOCW     45          0      One year              0
4  9237-HQITU      2          1  Month-to-month              1

          PaymentMethod  MonthlyCharges  TotalCharges  Churn  gender  \
0      Electronic check          29.85          29.85     0  Female
1          Mailed check          56.95         1889.5     0   Male
2          Mailed check          53.85          108.15     1   Male
3  Bank transfer (automatic)         42.30         1840.75     0   Male
4      Electronic check          70.70          151.65     1  Female

...          OnlineBackup_No  OnlineBackup_Yes  DeviceProtection_No  \

```

0	...	0	1	1
1	...	1	0	0
2	...	0	1	1
3	...	1	0	0
4	...	1	0	1

	DeviceProtection_Yes	TechSupport_No	TechSupport_Yes	StreamingTV_No	\
0	0	1	0	1	
1	1	1	0	1	
2	0	1	0	1	
3	1	0	1	1	
4	0	1	0	1	

	StreamingTV_Yes	StreamingMovies_No	StreamingMovies_Yes
0	0	1	0
1	0	1	0
2	0	1	0
3	0	1	0
4	0	1	0

[5 rows x 43 columns]

## Dropping the repeated variables

In [18]: *# We have created dummies for the below variables, so we can drop them*

```
telecom = telecom.drop(['Contract', 'PaymentMethod', 'gender', 'MultipleLines', 'InternetServiceProvider',
                        'TechSupport', 'StreamingTV', 'StreamingMovies'], 1)
```

In [19]: *#The variable was imported as a string we need to convert it to float*

```
telecom['TotalCharges'] = telecom['TotalCharges'].convert_objects(convert_numeric=True)
```

In [20]: telecom.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 7043 entries, 0 to 7042
```

```
Data columns (total 32 columns):
```

customerID	7043 non-null object
tenure	7043 non-null int64
PhoneService	7043 non-null int64
PaperlessBilling	7043 non-null int64
MonthlyCharges	7043 non-null float64
TotalCharges	7032 non-null float64
Churn	7043 non-null int64
SeniorCitizen	7043 non-null int64
Partner	7043 non-null int64
Dependents	7043 non-null int64
Contract_One year	7043 non-null uint8
Contract_Two year	7043 non-null uint8



```

PaymentMethod_Credit card (automatic)    7043 non-null uint8
PaymentMethod_Electronic check           7043 non-null uint8
PaymentMethod_Mailed check               7043 non-null uint8
gender_Male                             7043 non-null uint8
InternetService_Fiber optic              7043 non-null uint8
InternetService_No                       7043 non-null uint8
MultipleLines_No                         7043 non-null uint8
MultipleLines_Yes                       7043 non-null uint8
OnlineSecurity_No                       7043 non-null uint8
OnlineSecurity_Yes                      7043 non-null uint8
OnlineBackup_No                         7043 non-null uint8
OnlineBackup_Yes                       7043 non-null uint8
DeviceProtection_No                     7043 non-null uint8
DeviceProtection_Yes                    7043 non-null uint8
TechSupport_No                         7043 non-null uint8
TechSupport_Yes                       7043 non-null uint8
StreamingTV_No                         7043 non-null uint8
StreamingTV_Yes                       7043 non-null uint8
StreamingMovies_No                     7043 non-null uint8
StreamingMovies_Yes                    7043 non-null uint8
dtypes: float64(2), int64(7), object(1), uint8(22)
memory usage: 756.6+ KB

```

Now you can see that you have all variables as numeric.

## Checking for Outliers

```

In [21]: # Checking for outliers in the continuous variables
num_telecom = telecom[['tenure', 'MonthlyCharges', 'SeniorCitizen', 'TotalCharges']]

In [22]: # Checking outliers at 25%, 50%, 75%, 90%, 95% and 99%
num_telecom.describe(percentiles=[.25, .5, .75, .90, .95, .99])

Out[22]:

```

	tenure	MonthlyCharges	SeniorCitizen	TotalCharges
count	7043.000000	7043.000000	7043.000000	7032.000000
mean	32.371149	64.761692	0.162147	2283.300441
std	24.559481	30.090047	0.368612	2266.771362
min	0.000000	18.250000	0.000000	18.800000
25%	9.000000	35.500000	0.000000	401.450000
50%	29.000000	70.350000	0.000000	1397.475000
75%	55.000000	89.850000	0.000000	3794.737500
90%	69.000000	102.600000	1.000000	5976.640000
95%	72.000000	107.400000	1.000000	6923.590000
99%	72.000000	114.729000	1.000000	8039.883000
max	72.000000	118.750000	1.000000	8684.800000

From the distribution shown above, you can see that there no outliers in your data. The numbers are gradually increasing.

## Checking for Missing Values and Inputing Them

```
In [23]: # Adding up the missing values (column-wise)
         telecom.isnull().sum()
```

```
Out[23]: customerID          0
         tenure              0
         PhoneService        0
         PaperlessBilling    0
         MonthlyCharges      0
         TotalCharges        11
         Churn               0
         SeniorCitizen       0
         Partner             0
         Dependents          0
         Contract_One year   0
         Contract_Two year   0
         PaymentMethod_Credit card (automatic) 0
         PaymentMethod_Electronic check        0
         PaymentMethod_Mailed check            0
         gender_Male                0
         InternetService_Fiber optic          0
         InternetService_No              0
         MultipleLines_No                0
         MultipleLines_Yes              0
         OnlineSecurity_No              0
         OnlineSecurity_Yes            0
         OnlineBackup_No               0
         OnlineBackup_Yes              0
         DeviceProtection_No           0
         DeviceProtection_Yes          0
         TechSupport_No                0
         TechSupport_Yes              0
         StreamingTV_No                0
         StreamingTV_Yes              0
         StreamingMovies_No            0
         StreamingMovies_Yes           0
         dtype: int64
```

It means that  $11/7043 = 0.001561834$  i.e 0.1%, best is to remove these observations from the analysis

```
In [24]: # Checking the percentage of missing values
         round(100*(telecom.isnull().sum()/len(telecom.index)), 2)
```

```
Out[24]: customerID          0.00
         tenure              0.00
         PhoneService        0.00
         PaperlessBilling    0.00
```

MonthlyCharges	0.00
TotalCharges	0.16
Churn	0.00
SeniorCitizen	0.00
Partner	0.00
Dependents	0.00
Contract_One year	0.00
Contract_Two year	0.00
PaymentMethod_Credit card (automatic)	0.00
PaymentMethod_Electronic check	0.00
PaymentMethod_Mailed check	0.00
gender_Male	0.00
InternetService_Fiber optic	0.00
InternetService_No	0.00
MultipleLines_No	0.00
MultipleLines_Yes	0.00
OnlineSecurity_No	0.00
OnlineSecurity_Yes	0.00
OnlineBackup_No	0.00
OnlineBackup_Yes	0.00
DeviceProtection_No	0.00
DeviceProtection_Yes	0.00
TechSupport_No	0.00
TechSupport_Yes	0.00
StreamingTV_No	0.00
StreamingTV_Yes	0.00
StreamingMovies_No	0.00
StreamingMovies_Yes	0.00
dtype: float64	

```
In [25]: # Removing NaN TotalCharges rows
telecom = telecom[~np.isnan(telecom['TotalCharges'])]
```

```
In [26]: # Checking percentage of missing values after removing the missing values
round(100*(telecom.isnull().sum()/len(telecom.index)), 2)
```

```
Out[26]: customerID      0.0
tenure      0.0
PhoneService 0.0
PaperlessBilling 0.0
MonthlyCharges 0.0
TotalCharges 0.0
Churn      0.0
SeniorCitizen 0.0
Partner    0.0
Dependents 0.0
Contract_One year 0.0
Contract_Two year 0.0
```

```

PaymentMethod_Credit card (automatic)    0.0
PaymentMethod_Electronic check           0.0
PaymentMethod_Mailed check               0.0
gender_Male                             0.0
InternetService_Fiber optic              0.0
InternetService_No                       0.0
MultipleLines_No                         0.0
MultipleLines_Yes                       0.0
OnlineSecurity_No                       0.0
OnlineSecurity_Yes                      0.0
OnlineBackup_No                         0.0
OnlineBackup_Yes                       0.0
DeviceProtection_No                     0.0
DeviceProtection_Yes                    0.0
TechSupport_No                          0.0
TechSupport_Yes                         0.0
StreamingTV_No                          0.0
StreamingTV_Yes                         0.0
StreamingMovies_No                      0.0
StreamingMovies_Yes                     0.0
dtype: float64

```

Now we don't have any missing values

#### 0.1.4 Step 4: Test-Train Split

```
In [27]: from sklearn.model_selection import train_test_split
```

```
In [28]: # Putting feature variable to X
X = telecom.drop(['Churn', 'customerID'], axis=1)
```

```
X.head()
```

```

Out[28]:   tenure  PhoneService  PaperlessBilling  MonthlyCharges  TotalCharges  \
0         1             0             1         29.85         29.85
1        34             1             0         56.95        1889.50
2         2             1             1         53.85         108.15
3        45             0             0         42.30        1840.75
4         2             1             1         70.70         151.65

   SeniorCitizen  Partner  Dependents  Contract_One year  Contract_Two year  \
0              0        1           0              0              0
1              0        0           0              1              0
2              0        0           0              0              0
3              0        0           0              1              0
4              0        0           0              0              0

   ...  OnlineBackup_No  OnlineBackup_Yes  \
0    ...              0              1

```

1	...	1	0
2	...	0	1
3	...	1	0
4	...	1	0

	DeviceProtection_No	DeviceProtection_Yes	TechSupport_No	TechSupport_Yes	\
0	1	0	1	0	
1	0	1	1	0	
2	1	0	1	0	
3	0	1	0	1	
4	1	0	1	0	

	StreamingTV_No	StreamingTV_Yes	StreamingMovies_No	StreamingMovies_Yes
0	1	0	1	0
1	1	0	1	0
2	1	0	1	0
3	1	0	1	0
4	1	0	1	0

[5 rows x 30 columns]

```
In [29]: # Putting response variable to y
y = telecom['Churn']
```

```
y.head()
```

```
Out[29]: 0    0
1    0
2    1
3    0
4    1
Name: Churn, dtype: int64
```

```
In [30]: # Splitting the data into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, test_size=0.3)
```

### 0.1.5 Step 5: Feature Scaling

```
In [31]: from sklearn.preprocessing import StandardScaler
```

```
In [32]: scaler = StandardScaler()
```

```
X_train[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.fit_transform(X_train[['tenure', 'MonthlyCharges', 'TotalCharges']])
X_train.head()
```

```
Out[32]:      tenure  PhoneService  PaperlessBilling  MonthlyCharges  TotalCharges  \
879    0.019693             1             1      -0.338074      -0.276449
5790   0.305384             0             1      -0.464443      -0.112702
```

6498	-1.286319	1	1	0.581425	-0.974430
880	-0.919003	1	1	1.505913	-0.550676
2784	-1.163880	1	1	1.106854	-0.835971

	SeniorCitizen	Partner	Dependents	Contract_One year	\
879	0	0	0	0	
5790	0	1	1	0	
6498	0	0	0	0	
880	0	0	0	0	
2784	0	0	1	0	

	Contract_Two year	...	OnlineBackup_No	\
879	0	...	0	
5790	0	...	0	
6498	0	...	0	
880	0	...	0	
2784	0	...	1	

	OnlineBackup_Yes	DeviceProtection_No	DeviceProtection_Yes	\
879	1	1	0	
5790	1	1	0	
6498	1	0	1	
880	1	0	1	
2784	0	0	1	

	TechSupport_No	TechSupport_Yes	StreamingTV_No	StreamingTV_Yes	\
879	1	0	1	0	
5790	1	0	0	1	
6498	1	0	1	0	
880	0	1	0	1	
2784	0	1	0	1	

	StreamingMovies_No	StreamingMovies_Yes
879	1	0
5790	0	1
6498	1	0
880	0	1
2784	0	1

[5 rows x 30 columns]

```
In [33]: ### Checking the Churn Rate
churn = (sum(telecom['Churn'])/len(telecom['Churn'].index))*100
churn
```


Out[33]: 26.578498293515356

We have almost 27% churn rate

## 0.1.6 Step 6: Looking at Correlations

```
In [34]: # Importing matplotlib and seaborn
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [35]: # Let's see the correlation matrix
plt.figure(figsize = (20,10)) # Size of the figure
sns.heatmap(telecom.corr(),annot = True)
plt.show()
```




## Dropping highly correlated dummy variables

```
In [36]: X_test = X_test.drop(['MultipleLines_No', 'OnlineSecurity_No', 'OnlineBackup_No', 'DeviceProtection_No', 'TechSupport_No', 'StreamingTV_No', 'StreamingMovies_No'], 1)
X_train = X_train.drop(['MultipleLines_No', 'OnlineSecurity_No', 'OnlineBackup_No', 'DeviceProtection_No', 'TechSupport_No', 'StreamingTV_No', 'StreamingMovies_No'], 1)
```

**Checking the Correlation Matrix** After dropping highly correlated variables now let's check the correlation matrix again.

```
In [37]: plt.figure(figsize = (20,10))
sns.heatmap(X_train.corr(),annot = True)
plt.show()
```



## 0.1.7 Step 7: Model Building

Let's start by splitting our data into a training set and a test set.

### Running Your First Training Model

```
In [38]: import statsmodels.api as sm
```

```
In [39]: # Logistic regression model
```

```
logm1 = sm.GLM(y_train,(sm.add_constant(X_train))), family = sm.families.Binomial())
logm1.fit().summary()
```

```
Out[39]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

#### Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          4922
Model:                GLM      Df Residuals:              4898
Model Family:         Binomial  Df Model:                  23
Link Function:         logit    Scale:                    1.0000
Method:                IRLS     Log-Likelihood:          -2004.7
Date:                 Thu, 29 Nov 2018    Deviance:                 4009.4
Time:                 11:23:01    Pearson chi2:            6.07e+03
No. Iterations:        7         Covariance Type:         nonrobust
=====
```



	coef	std err	z	P> z
-----				
const	-3.9382	1.546	-2.547	0.011
tenure	-1.5172	0.189	-8.015	0.000
PhoneService	0.9507	0.789	1.205	0.228
PaperlessBilling	0.3254	0.090	3.614	0.000
MonthlyCharges	-2.1806	1.160	-1.880	0.060
TotalCharges	0.7332	0.198	3.705	0.000
SeniorCitizen	0.3984	0.102	3.924	0.000
Partner	0.0374	0.094	0.399	0.690
Dependents	-0.1430	0.107	-1.332	0.183
Contract_One year	-0.6578	0.129	-5.106	0.000
Contract_Two year	-1.2455	0.212	-5.874	0.000
PaymentMethod_Credit card (automatic)	-0.2577	0.137	-1.883	0.060
PaymentMethod_Electronic check	0.1615	0.113	1.434	0.152
PaymentMethod_Mailed check	-0.2536	0.137	-1.845	0.065
gender_Male	-0.0346	0.078	-0.442	0.658
InternetService_Fiber optic	2.5124	0.967	2.599	0.009
InternetService_No	-2.7792	0.982	-2.831	0.005
MultipleLines_Yes	0.5623	0.214	2.628	0.009
OnlineSecurity_Yes	-0.0245	0.216	-0.113	0.910
OnlineBackup_Yes	0.1740	0.212	0.822	0.411
DeviceProtection_Yes	0.3229	0.215	1.501	0.133
TechSupport_Yes	-0.0305	0.216	-0.141	0.888
StreamingTV_Yes	0.9598	0.396	2.423	0.015
StreamingMovies_Yes	0.8484	0.396	2.143	0.032
=====				
"""				

### 0.1.8 Step 8: Feature Selection Using RFE

```
In [40]: from sklearn.linear_model import LogisticRegression
         logreg = LogisticRegression()

In [41]: from sklearn.feature_selection import RFE
         rfe = RFE(logreg, 15)           # running RFE with 13 variables as output
         rfe = rfe.fit(X_train, y_train)

In [42]: rfe.support_

Out[42]: array([ True,  True,  True, False,  True,  True, False, False,  True,
                True,  True, False,  True, False,  True,  True,  True,  True,
                False, False,  True,  True, False])

In [43]: list(zip(X_train.columns, rfe.support_, rfe.ranking_))

Out[43]: [('tenure', True, 1),
          ('PhoneService', True, 1),
          ('PaperlessBilling', True, 1),
```

```
( 'MonthlyCharges', False, 6),
( 'TotalCharges', True, 1),
( 'SeniorCitizen', True, 1),
( 'Partner', False, 8),
( 'Dependents', False, 4),
( 'Contract_One year', True, 1),
( 'Contract_Two year', True, 1),
( 'PaymentMethod_Credit card (automatic)', True, 1),
( 'PaymentMethod_Electronic check', False, 3),
( 'PaymentMethod_Mailed check', True, 1),
( 'gender_Male', False, 9),
( 'InternetService_Fiber optic', True, 1),
( 'InternetService_No', True, 1),
( 'MultipleLines_Yes', True, 1),
( 'OnlineSecurity_Yes', True, 1),
( 'OnlineBackup_Yes', False, 2),
( 'DeviceProtection_Yes', False, 7),
( 'TechSupport_Yes', True, 1),
( 'StreamingTV_Yes', True, 1),
( 'StreamingMovies_Yes', False, 5)]
```

```
In [44]: col = X_train.columns[rfe.support_]
```

```
In [45]: X_train.columns[~rfe.support_]
```

```
Out [45]: Index(['MonthlyCharges', 'Partner', 'Dependents',
                'PaymentMethod_Electronic check', 'gender_Male', 'OnlineBackup_Yes',
                'DeviceProtection_Yes', 'StreamingMovies_Yes'],
                dtype='object')
```

## Assessing the model with StatsModels

```
In [46]: X_train_sm = sm.add_constant(X_train[col])
logm2 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
res = logm2.fit()
res.summary()
```

```
Out [46]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                Generalized Linear Model Regression Results
=====
Dep. Variable:                  Churn    No. Observations:                  4922
Model:                          GLM      Df Residuals:                      4906
Model Family:                   Binomial  Df Model:                          15
Link Function:                  logit     Scale:                            1.0000
Method:                         IRLS     Log-Likelihood:                    -2011.8
Date:                           Thu, 29 Nov 2018  Deviance:                     4023.5
Time:                           11:23:04  Pearson chi2:                      6.22e+03
No. Iterations:                  7        Covariance Type:                  nonrobust
```

	coef	std err	z	P> z
const	-1.0343	0.171	-6.053	0.000
tenure	-1.5386	0.184	-8.381	0.000
PhoneService	-0.5231	0.161	-3.256	0.001
PaperlessBilling	0.3397	0.090	3.789	0.000
TotalCharges	0.7116	0.188	3.794	0.000
SeniorCitizen	0.4294	0.100	4.312	0.000
Contract_One year	-0.6813	0.128	-5.334	0.000
Contract_Two year	-1.2680	0.211	-6.011	0.000
PaymentMethod_Credit card (automatic)	-0.3775	0.113	-3.352	0.001
PaymentMethod_Mailed check	-0.3760	0.111	-3.389	0.001
InternetService_Fiber optic	0.7421	0.117	6.317	0.000
InternetService_No	-0.9385	0.166	-5.650	0.000
MultipleLines_Yes	0.2086	0.096	2.181	0.029
OnlineSecurity_Yes	-0.4049	0.102	-3.968	0.000
TechSupport_Yes	-0.3967	0.102	-3.902	0.000
StreamingTV_Yes	0.2747	0.094	2.911	0.004

In [47]: # Getting the predicted values on the train set

```
y_train_pred = res.predict(X_train_sm)
y_train_pred[:10]
```

```
Out[47]: 879      0.225111
5790     0.274893
6498     0.692126
880      0.504909
2784     0.645261
3874     0.417544
5387     0.420131
6623     0.809427
4465     0.223211
5364     0.512246
dtype: float64
```

```
In [48]: y_train_pred = y_train_pred.values.reshape(-1)
y_train_pred[:10]
```

```
Out[48]: array([0.22511138, 0.27489289, 0.69212611, 0.50490896, 0.6452606 ,
                0.41754449, 0.42013086, 0.80942651, 0.2232105 , 0.51224637])
```

### Creating a dataframe with the actual churn flag and the predicted probabilities

```
In [49]: y_train_pred_final = pd.DataFrame({'Churn':y_train.values, 'Churn_Prob':y_train_pred})
y_train_pred_final['CustID'] = y_train.index
y_train_pred_final.head()
```

```
Out [49]:
```

	Churn	Churn_Prob	CustID
0	0	0.225111	879
1	0	0.274893	5790
2	1	0.692126	6498
3	1	0.504909	880
4	1	0.645261	2784

Creating new column 'predicted' with 1 if Churn\_Prob > 0.5 else 0

```
In [50]: y_train_pred_final['predicted'] = y_train_pred_final.Churn_Prob.map(lambda x: 1 if x > 0.5 else 0)

# Let's see the head
y_train_pred_final.head()
```

```
Out [50]:
```

	Churn	Churn_Prob	CustID	predicted
0	0	0.225111	879	0
1	0	0.274893	5790	0
2	1	0.692126	6498	1
3	1	0.504909	880	1
4	1	0.645261	2784	1

```
In [51]: from sklearn import metrics
```

```
In [52]: # Confusion matrix
confusion = metrics.confusion_matrix(y_train_pred_final.Churn, y_train_pred_final.predicted)
print(confusion)
```

```
[[3270  365]
 [ 579  708]]
```

```
In [53]: # Predicted      not_churn      churn
# Actual
# not_churn      3270      365
# churn          579      708
```

```
In [54]: # Let's check the overall accuracy.
print(metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.predicted))
```

```
0.8082080455099553
```

## Checking VIFs

```
In [55]: # Check for the VIF values of the feature variables.
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [56]: # Create a dataframe that will contain the names of all the feature variables and the
vif = pd.DataFrame()
```

```

vif['Features'] = X_train[col].columns
vif['VIF'] = [variance_inflation_factor(X_train[col].values, i) for i in range(X_train[
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

```

Out [56]:
           Features  VIF
1      PhoneService  8.86
3      TotalCharges  7.37
0           tenure  6.88
9  InternetService_Fiber optic  3.97
6      Contract_Two year  3.28
10  InternetService_No  3.25
2      PaperlessBilling  2.68
11  MultipleLines_Yes  2.53
14  StreamingTV_Yes  2.34
13  TechSupport_Yes  2.08
5      Contract_One year  1.93
12  OnlineSecurity_Yes  1.90
8  PaymentMethod_Mailed check  1.72
7  PaymentMethod_Credit card (automatic)  1.46
4      SeniorCitizen  1.31

```

There are a few variables with high VIF. It's best to drop these variables as they aren't helping much with prediction and unnecessarily making the model complex. The variable 'PhoneService' has the highest VIF. So let's start by dropping that.

```

In [57]: col = col.drop('PhoneService', 1)
col

```

```

Out [57]: Index(['tenure', 'PaperlessBilling', 'TotalCharges', 'SeniorCitizen',
                'Contract_One year', 'Contract_Two year',
                'PaymentMethod_Credit card (automatic)', 'PaymentMethod_Mailed check',
                'InternetService_Fiber optic', 'InternetService_No',
                'MultipleLines_Yes', 'OnlineSecurity_Yes', 'TechSupport_Yes',
                'StreamingTV_Yes'],
                dtype='object')

```

```

In [58]: # Let's re-run the model using the selected variables
X_train_sm = sm.add_constant(X_train[col])
logm3 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
res = logm3.fit()
res.summary()

```

```

Out [58]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                Generalized Linear Model Regression Results
=====
Dep. Variable:                  Churn    No. Observations:                  4922

```

```

Model: GLM Df Residuals: 4907
Model Family: Binomial Df Model: 14
Link Function: logit Scale: 1.0000
Method: IRLS Log-Likelihood: -2017.0
Date: Thu, 29 Nov 2018 Deviance: 4034.0
Time: 11:23:05 Pearson chi2: 5.94e+03
No. Iterations: 7 Covariance Type: nonrobust

```

```

=====
              coef      std err          z      P>|z|
-----
const          -1.3885      0.133    -10.437      0.000
tenure          -1.4138      0.179     -7.884      0.000
PaperlessBilling    0.3425      0.089      3.829      0.000
TotalCharges       0.5936      0.184      3.225      0.001
SeniorCitizen      0.4457      0.099      4.486      0.000
Contract_One year  -0.6905      0.128     -5.411      0.000
Contract_Two year  -1.2646      0.211     -6.002      0.000
PaymentMethod_Credit card (automatic) -0.3785      0.113     -3.363      0.001
PaymentMethod_Mailed check -0.3769      0.111     -3.407      0.001
InternetService_Fiber optic    0.6241      0.111      5.645      0.000
InternetService_No  -1.0940      0.158     -6.919      0.000
MultipleLines_Yes    0.1607      0.094      1.712      0.087
OnlineSecurity_Yes  -0.4094      0.102     -4.016      0.000
TechSupport_Yes     -0.4085      0.101     -4.025      0.000
StreamingTV_Yes      0.3077      0.094      3.277      0.001
=====
"""

```

```
In [59]: y_train_pred = res.predict(X_train_sm).values.reshape(-1)
```

```
In [60]: y_train_pred[:10]
```

```
Out[60]: array([0.25403236, 0.22497676, 0.69386521, 0.51008735, 0.65172434,
                0.45441958, 0.3272777 , 0.80583357, 0.17618503, 0.50403034])
```

```
In [61]: y_train_pred_final['Churn_Prob'] = y_train_pred
```

```
In [62]: # Creating new column 'predicted' with 1 if Churn_Prob > 0.5 else 0
```

```

y_train_pred_final['predicted'] = y_train_pred_final.Churn_Prob.map(lambda x: 1 if x > 0.5 else 0)
y_train_pred_final.head()

```

```

Out[62]:   Churn  Churn_Prob  CustID  predicted
0      0      0.254032      879         0
1      0      0.224977      5790        0
2      1      0.693865      6498         1
3      1      0.510087       880         1
4      1      0.651724      2784         1

```

```
In [63]: # Let's check the overall accuracy.
```

```
print(metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.predicted))
```

0.8051605038602194

So overall the accuracy hasn't dropped much.

### Let's check the VIFs again

```
In [64]: vif = pd.DataFrame()
vif['Features'] = X_train[col].columns
vif['VIF'] = [variance_inflation_factor(X_train[col].values, i) for i in range(X_train[col].shape[0])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
Out [64]:
```

	Features	VIF
2	TotalCharges	7.30
0	tenure	6.79
5	Contract_Two year	3.16
8	InternetService_Fiber optic	2.94
9	InternetService_No	2.53
1	PaperlessBilling	2.52
13	StreamingTV_Yes	2.31
10	MultipleLines_Yes	2.27
12	TechSupport_Yes	2.00
4	Contract_One year	1.83
11	OnlineSecurity_Yes	1.80
7	PaymentMethod_Mailed check	1.66
6	PaymentMethod_Credit card (automatic)	1.44
3	SeniorCitizen	1.31

```
In [65]: # Let's drop TotalCharges since it has a high VIF
col = col.drop('TotalCharges')
col
```

```
Out [65]: Index(['tenure', 'PaperlessBilling', 'SeniorCitizen', 'Contract_One year',
                 'Contract_Two year', 'PaymentMethod_Credit card (automatic)',
                 'PaymentMethod_Mailed check', 'InternetService_Fiber optic',
                 'InternetService_No', 'MultipleLines_Yes', 'OnlineSecurity_Yes',
                 'TechSupport_Yes', 'StreamingTV_Yes'],
                 dtype='object')
```

```
In [66]: # Let's re-run the model using the selected variables
X_train_sm = sm.add_constant(X_train[col])
logm4 = sm.GLM(y_train, X_train_sm, family = sm.families.Binomial())
res = logm4.fit()
res.summary()
```

```
Out [66]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

# Generalized Linear Model Regression Results

```

=====
Dep. Variable:          Churn    No. Observations:          4922
Model:                  GLM      Df Residuals:              4908
Model Family:          Binomial  Df Model:                  13
Link Function:          logit    Scale:                    1.0000
Method:                 IRLS     Log-Likelihood:           -2022.5
Date:                   Thu, 29 Nov 2018    Deviance:                  4044.9
Time:                   11:23:06    Pearson chi2:              5.22e+03
No. Iterations:         7         Covariance Type:          nonrobust
=====

```

	coef	std err	z	P> z
const	-1.4695	0.130	-11.336	0.000
tenure	-0.8857	0.065	-13.553	0.000
PaperlessBilling	0.3367	0.089	3.770	0.000
SeniorCitizen	0.4517	0.100	4.527	0.000
Contract_One year	-0.6792	0.127	-5.360	0.000
Contract_Two year	-1.2308	0.208	-5.903	0.000
PaymentMethod_Credit card (automatic)	-0.3827	0.113	-3.399	0.001
PaymentMethod_Mailed check	-0.3393	0.110	-3.094	0.002
InternetService_Fiber optic	0.7914	0.098	8.109	0.000
InternetService_No	-1.1205	0.157	-7.127	0.000
MultipleLines_Yes	0.2166	0.092	2.355	0.019
OnlineSecurity_Yes	-0.3739	0.101	-3.684	0.000
TechSupport_Yes	-0.3611	0.101	-3.591	0.000
StreamingTV_Yes	0.3995	0.089	4.465	0.000

"""

```
In [67]: y_train_pred = res.predict(X_train_sm).values.reshape(-1)
```

```
In [68]: y_train_pred[:10]
```

```
Out [68]: array([0.28219274, 0.2681923 , 0.68953115, 0.53421409, 0.67433213,
0.42980951, 0.31009304, 0.81248467, 0.20462744, 0.50431479])
```

```
In [69]: y_train_pred_final['Churn_Prob'] = y_train_pred
```

```
In [70]: # Creating new column 'predicted' with 1 if Churn_Prob > 0.5 else 0
```

```
y_train_pred_final['predicted'] = y_train_pred_final.Churn_Prob.map(lambda x: 1 if x > 0.5 else 0)
y_train_pred_final.head()
```

```
Out [70]:
```

	Churn	Churn_Prob	CustID	predicted
0	0	0.282193	879	0
1	0	0.268192	5790	0
2	1	0.689531	6498	1
3	1	0.534214	880	1
4	1	0.674332	2784	1



```
In [71]: # Let's check the overall accuracy.
         print(metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.predicted))

0.804754164973588
```

The accuracy is still practically the same.

**Let's now check the VIFs again**

```
In [72]: vif = pd.DataFrame()
         vif['Features'] = X_train[col].columns
         vif['VIF'] = [variance_inflation_factor(X_train[col].values, i) for i in range(X_train[
         vif['VIF'] = round(vif['VIF'], 2)
         vif = vif.sort_values(by = "VIF", ascending = False)
         vif
```

```
Out[72]:
```

	Features	VIF
4	Contract_Two year	3.07
7	InternetService_Fiber optic	2.60
1	PaperlessBilling	2.44
9	MultipleLines_Yes	2.24
12	StreamingTV_Yes	2.17
8	InternetService_No	2.12
0	tenure	2.04
11	TechSupport_Yes	1.98
3	Contract_One year	1.82
10	OnlineSecurity_Yes	1.78
6	PaymentMethod_Mailed check	1.66
5	PaymentMethod_Credit card (automatic)	1.44
2	SeniorCitizen	1.31

All variables have a good value of VIF. So we need not drop any more variables and we can proceed with making predictions using this model only

```
In [73]: # Let's take a look at the confusion matrix again
         confusion = metrics.confusion_matrix(y_train_pred_final.Churn, y_train_pred_final.pre
         confusion
```

```
Out[73]: array([[3269,  366],
                [ 595,  692]], dtype=int64)
```

```
In [74]: # Actual/Predicted      not_churn      churn
         # not_churn              3269          366
         # churn                  595           692
```

```
In [75]: # Let's check the overall accuracy.
         metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.predicted)
```

```
Out[75]: 0.804754164973588
```

## 0.2 Metrics beyond simply accuracy

```
In [76]: TP = confusion[1,1] # true positive
         TN = confusion[0,0] # true negatives
         FP = confusion[0,1] # false positives
         FN = confusion[1,0] # false negatives
```

```
In [77]: # Let's see the sensitivity of our logistic regression model
         TP / float(TP+FN)
```

```
Out[77]: 0.5376845376845377
```

```
In [78]: # Let us calculate specificity
         TN / float(TN+FP)
```

```
Out[78]: 0.8993122420907841
```

```
In [79]: # Calculate false positive rate - predicting churn when customer does not have churned
         print(FP/ float(TN+FP))
```

```
0.10068775790921596
```

```
In [80]: # positive predictive value
         print (TP / float(TP+FP))
```

```
0.6540642722117203
```

```
In [81]: # Negative predictive value
         print (TN / float(TN+ FN))
```

```
0.8460144927536232
```

### 0.2.1 Step 9: Plotting the ROC Curve

An ROC curve demonstrates several things:

- It shows the tradeoff between sensitivity and specificity (any increase in sensitivity will be accompanied by a decrease in specificity).
- The closer the curve follows the left-hand border and then the top border of the ROC space, the more accurate the test.
- The closer the curve comes to the 45-degree diagonal of the ROC space, the less accurate the test.

```
In [82]: def draw_roc( actual, probs ):
         fpr, tpr, thresholds = metrics.roc_curve( actual, probs,
                                                    drop_intermediate = False )
         auc_score = metrics.roc_auc_score( actual, probs )
```

```

plt.figure(figsize=(5, 5))
plt.plot( fpr, tpr, label='ROC curve (area = %0.2f)' % auc_score )
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

```

```

return None

```

```


In [83]: fpr, tpr, thresholds = metrics.roc_curve( y_train_pred_final.Churn, y_train_pred_final.Churn_Prob)

```

```

In [84]: draw_roc(y_train_pred_final.Churn, y_train_pred_final.Churn_Prob)

```



## 0.2.2 Step 10: Finding Optimal Cutoff Point

Optimal cutoff probability is that prob where we get balanced sensitivity and specificity

```
In [85]: # Let's create columns with different probability cutoffs
numbers = [float(x)/10 for x in range(10)]
for i in numbers:
    y_train_pred_final[i]= y_train_pred_final.Churn_Prob.map(lambda x: 1 if x > i else 0)
y_train_pred_final.head()
```

```
Out [85]:
```

	Churn	Churn_Prob	CustID	predicted	0.0	0.1	0.2	0.3	0.4	0.5	0.6	\
0	0	0.282193	879	0	1	1	1	0	0	0	0	
1	0	0.268192	5790	0	1	1	1	0	0	0	0	
2	1	0.689531	6498	1	1	1	1	1	1	1	1	
3	1	0.534214	880	1	1	1	1	1	1	1	0	
4	1	0.674332	2784	1	1	1	1	1	1	1	1	

	0.7	0.8	0.9
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

```
In [86]: # Now let's calculate accuracy sensitivity and specificity for various probability cu
cutoff_df = pd.DataFrame( columns = ['prob', 'accuracy', 'sensi', 'speci'])
from sklearn.metrics import confusion_matrix

# TP = confusion[1,1] # true positive
# TN = confusion[0,0] # true negatives
# FP = confusion[0,1] # false positives
# FN = confusion[1,0] # false negatives


num = [0.0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9]
for i in num:
    cm1 = metrics.confusion_matrix(y_train_pred_final.Churn, y_train_pred_final[i] )
    total1=sum(sum(cm1))
    accuracy = (cm1[0,0]+cm1[1,1])/total1

    speci = cm1[0,0]/(cm1[0,0]+cm1[0,1])
    sensi = cm1[1,1]/(cm1[1,0]+cm1[1,1])
    cutoff_df.loc[i] =[ i ,accuracy,sensi,speci]
print(cutoff_df)
```

	prob	accuracy	sensi	speci
0.0	0.0	0.261479	1.000000	0.000000
0.1	0.1	0.619667	0.946387	0.503989
0.2	0.2	0.722674	0.850039	0.677579
0.3	0.3	0.771434	0.780109	0.768363
0.4	0.4	0.795002	0.671329	0.838790
0.5	0.5	0.804754	0.537685	0.899312
0.6	0.6	0.800284	0.385392	0.947180

0.7	0.7	0.779764	0.205128	0.983219
0.8	0.8	0.749289	0.050505	0.996699
0.9	0.9	0.738521	0.000000	1.000000

```
In [87]: # Let's plot accuracy sensitivity and specificity for various probabilities.
cutoff_df.plot.line(x='prob', y=['accuracy', 'sensi', 'speci'])
plt.show()
```



From the curve above, 0.3 is the optimum point to take it as a cutoff probability.

```
In [88]: y_train_pred_final['final_predicted'] = y_train_pred_final.Churn_Prob.map( lambda x:
y_train_pred_final.head()
```

```
Out[88]:
```

	Churn	Churn_Prob	CustID	predicted	0.0	0.1	0.2	0.3	0.4	0.5	0.6	\
0	0	0.282193	879	0	1	1	1	0	0	0	0	
1	0	0.268192	5790	0	1	1	1	0	0	0	0	
2	1	0.689531	6498	1	1	1	1	1	1	1	1	
3	1	0.534214	880	1	1	1	1	1	1	1	0	
4	1	0.674332	2784	1	1	1	1	1	1	1	1	

	0.7	0.8	0.9	final_predicted
0	0	0	0	0
1	0	0	0	0

2	0	0	0	1
3	0	0	0	1
4	0	0	0	1

```
In [89]: # Let's check the overall accuracy.
metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.final_predicted)
```

```
Out[89]: 0.771434376269809
```

```
In [90]: confusion2 = metrics.confusion_matrix(y_train_pred_final.Churn, y_train_pred_final.final_predicted)
confusion2
```

```
Out[90]: array([[2793, 842],
               [ 283, 1004]], dtype=int64)
```

```
In [91]: TP = confusion2[1,1] # true positive
TN = confusion2[0,0] # true negatives
FP = confusion2[0,1] # false positives
FN = confusion2[1,0] # false negatives
```

```
In [92]: # Let's see the sensitivity of our logistic regression model
TP / float(TP+FN)
```

```
Out[92]: 0.7801087801087802
```

```
In [93]: # Let us calculate specificity
TN / float(TN+FP)
```

```
Out[93]: 0.768363136176066
```

```
In [94]: # Calculate false positive rate - predicting churn when customer does not have churned
print(FP/ float(TN+FP))
```

```
0.23163686382393398
```

```
In [95]: # Positive predictive value
print (TP / float(TP+FP))
```

```
0.5438786565547129
```

```
In [96]: # Negative predictive value
print (TN / float(TN+ FN))
```

```
0.907997399219766
```

### 0.3 Precision and Recall

```
In [97]: #Looking at the confusion matrix again
```

```
In [98]: confusion = metrics.confusion_matrix(y_train_pred_final.Churn, y_train_pred_final.predicted)
          confusion
```

```
Out[98]: array([[3269,  366],
                [ 595,  692]], dtype=int64)
```

**Precision**  $TP / TP + FP$

```
In [99]: confusion[1,1]/(confusion[0,1]+confusion[1,1])
```

```
Out[99]: 0.6540642722117203
```

**Recall**  $TP / TP + FN$

```
In [100]: confusion[1,1]/(confusion[1,0]+confusion[1,1])
```

```
Out[100]: 0.5376845376845377
```

Using sklearn utilities for the same

```
In [101]: from sklearn.metrics import precision_score, recall_score
```

```
In [102]: ?precision_score
```

```
In [103]: precision_score(y_train_pred_final.Churn, y_train_pred_final.predicted)
```

```
Out[103]: 0.6540642722117203
```

```
In [104]: recall_score(y_train_pred_final.Churn, y_train_pred_final.predicted)
```

```
Out[104]: 0.5376845376845377
```

#### 0.3.1 Precision and recall tradeoff

```
In [105]: from sklearn.metrics import precision_recall_curve
```

```
In [106]: y_train_pred_final.Churn, y_train_pred_final.predicted
```

```
Out[106]: (0      0
           1      0
           2      1
           3      1
           4      1
           5      0
           6      0
           7      1
           8      0)
```

9	1
10	0
11	1
12	1
13	0
14	0
15	0
16	0
17	0
18	0
19	0
20	0
21	0
22	0
23	0
24	0
25	0
26	0
27	0
28	0
29	0
	..
4892	1
4893	1
4894	0
4895	0
4896	0
4897	0
4898	0
4899	0
4900	0
4901	1
4902	0
4903	1
4904	0
4905	0
4906	1
4907	0
4908	0
4909	1
4910	0
4911	0
4912	0
4913	0
4914	0
4915	0
4916	1
4917	0




4918	0
4919	0
4920	0
4921	0
Name: Churn, Length: 4922, dtype: int64, 0	
1	0
2	1
3	1
4	1
5	0
6	0
7	1
8	0
9	1
10	0
11	1
12	1
13	0
14	0
15	0
16	0
17	0
18	0
19	0
20	0
21	0
22	0
23	0
24	0
25	0
26	0
27	0
28	0
29	0
..	
4892	0
4893	1
4894	0
4895	0
4896	0
4897	0
4898	0
4899	0
4900	0
4901	0
4902	0
4903	0
4904	1

```
4905    0
4906    1
4907    0
4908    0
4909    1
4910    0
4911    0
4912    0
4913    0
4914    0
4915    0
4916    0
4917    0
4918    0
4919    0
4920    0
4921    0
Name: predicted, Length: 4922, dtype: int64)
```

```
In [107]: p, r, thresholds = precision_recall_curve(y_train_pred_final.Churn, y_train_pred_final.Churn)
```

```
In [108]: plt.plot(thresholds, p[:-1], "g-")
plt.plot(thresholds, r[:-1], "r-")
plt.show()
```



### 0.3.2 Step 11: Making predictions on the test set

```
In [109]: X_test[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.transform(X_test[['tenure', 'MonthlyCharges', 'TotalCharges']])
```

```
In [110]: X_test = X_test[col]
X_test.head()
```

```
Out[110]:
```

	tenure	PaperlessBilling	SeniorCitizen	Contract_One year	\
942	-0.347623	1	0	0	
3730	0.999203	1	0	0	
1761	1.040015	1	0	0	
2283	-1.286319	1	0	0	
1872	0.346196	0	0	0	

	Contract_Two year	PaymentMethod_Credit card (automatic)	\
942	0	1	
3730	0	1	
1761	1	1	
2283	0	0	
1872	1	0	

	PaymentMethod_Mailed check	InternetService_Fiber optic	\
942	0	1	
3730	0	1	
1761	0	0	
2283	1	1	
1872	0	0	

	InternetService_No	MultipleLines_Yes	OnlineSecurity_Yes	\
942	0	0	0	
3730	0	1	0	
1761	1	1	0	
2283	0	0	0	
1872	1	0	0	

	TechSupport_Yes	StreamingTV_Yes
942	0	0
3730	0	1
1761	0	0
2283	0	0
1872	0	0

```
In [111]: X_test_sm = sm.add_constant(X_test)
```

Making predictions on the test set

```
In [112]: y_test_pred = res.predict(X_test_sm)
```

```
In [113]: y_test_pred[:10]
```

```
Out[113]: 942      0.397413
          3730     0.270295
          1761     0.010238
          2283     0.612692
          1872     0.015869
          1970     0.727206
          2532     0.302131
          1616     0.010315
          2485     0.632881
          5914     0.126451
          dtype: float64
```

```
In [114]: # Converting y_pred to a dataframe which is an array
          y_pred_1 = pd.DataFrame(y_test_pred)
```

```
In [115]: # Let's see the head
          y_pred_1.head()
```

```
Out[115]:      0
          942    0.397413
          3730    0.270295
          1761    0.010238
          2283    0.612692
          1872    0.015869
```

```
In [116]: # Converting y_test to dataframe
          y_test_df = pd.DataFrame(y_test)
```

```
In [117]: # Putting CustID to index
          y_test_df['CustID'] = y_test_df.index
```

```
In [118]: # Removing index for both dataframes to append them side by side
          y_pred_1.reset_index(drop=True, inplace=True)
          y_test_df.reset_index(drop=True, inplace=True)
```

```
In [119]: # Appending y_test_df and y_pred_1
          y_pred_final = pd.concat([y_test_df, y_pred_1],axis=1)
```

```
In [120]: y_pred_final.head()
```

```
Out[120]:   Churn  CustID      0
          0      942  0.397413
          1      3730  0.270295
          2       1761  0.010238
          3       2283  0.612692
          4       1872  0.015869
```

```
In [121]: # Renaming the column
          y_pred_final= y_pred_final.rename(columns={ 0 : 'Churn_Prob'})
```

```

In [122]: # Rearranging the columns
          y_pred_final = y_pred_final.reindex_axis(['CustID', 'Churn', 'Churn_Prob'], axis=1)

In [123]: # Let's see the head of y_pred_final
          y_pred_final.head()

Out[123]:
   CustID  Churn  Churn_Prob
0      942     0    0.397413
1     3730     1    0.270295
2     1761     0    0.010238
3     2283     1    0.612692
4     1872     0    0.015869

In [124]: y_pred_final['final_predicted'] = y_pred_final.Churn_Prob.map(lambda x: 1 if x > 0.4

In [125]: y_pred_final.head()

Out[125]:
   CustID  Churn  Churn_Prob  final_predicted
0      942     0    0.397413                0
1     3730     1    0.270295                0
2     1761     0    0.010238                0
3     2283     1    0.612692                1
4     1872     0    0.015869                0

In [126]: # Let's check the overall accuracy.
          metrics.accuracy_score(y_pred_final.Churn, y_pred_final.final_predicted)

Out[126]: 0.7834123222748816

In [127]: confusion2 = metrics.confusion_matrix(y_pred_final.Churn, y_pred_final.final_predict
          confusion2

Out[127]: array([[1294,  234],
                [ 223,  359]], dtype=int64)

In [128]: TP = confusion2[1,1] # true positive
          TN = confusion2[0,0] # true negatives
          FP = confusion2[0,1] # false positives
          FN = confusion2[1,0] # false negatives

In [129]: # Let's see the sensitivity of our logistic regression model
          TP / float(TP+FN)

Out[129]: 0.6168384879725086

In [130]: # Let us calculate specificity
          TN / float(TN+FP)

Out[130]: 0.8468586387434555

```