

SVM - Letter Recognition

October 22, 2020

1 Letter Recognition Using SVM

Let's now tackle a slightly more complex problem - letter recognition. We'll first explore the dataset a bit, prepare it (scale etc.) and then experiment with linear and non-linear SVMs with various hyperparameters.

1.1 Data Understanding

Let's first understand the shape, attributes etc. of the dataset.

```
In [1]: # libraries
import pandas as pd
import numpy as np
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import scale

# dataset
letters = pd.read_csv("letter-recognition.csv")

In [2]: # about the dataset

# dimensions
print("Dimensions: ", letters.shape, "\n")

# data types
print(letters.info())

# head
letters.head()
```

Dimensions: (20000, 17)

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 17 columns):
letter      20000 non-null object
xbox        20000 non-null int64
ybox        20000 non-null int64
width       20000 non-null int64
height      20000 non-null int64
onpix       20000 non-null int64
xbar        20000 non-null int64
ybar        20000 non-null int64
x2bar       20000 non-null int64
y2bar       20000 non-null int64
xybar       20000 non-null int64
x2ybar      20000 non-null int64
xy2bar      20000 non-null int64
xedge       20000 non-null int64
xedgey      20000 non-null int64
yedge       20000 non-null int64
yedgey      20000 non-null int64
dtypes: int64(16), object(1)
memory usage: 2.6+ MB
None
```

```
Out[2]:
```

	letter	xbox	ybox	width	height	onpix	xbar	ybar	x2bar	y2bar	\
0	T	2	8	3	5	1	8	13	0	6	
1	I	5	12	3	7	2	10	5	5	4	
2	D	4	11	6	8	6	10	6	2	6	
3	N	7	11	6	6	3	5	9	4	6	
4	G	2	1	3	1	1	8	6	6	6	

	xybar	x2ybar	xy2bar	xedge	xedgey	yedge	yedgey
0	6	10	8	0	8	0	8
1	13	3	9	2	8	4	10
2	10	3	7	3	7	3	9
3	4	4	10	6	10	2	8
4	6	5	9	1	7	5	10

```
In [3]: # a quirky bug: the column names have a space, e.g. 'xbox ', which throws an error when
print(letters.columns)
```

```
Index(['letter', 'xbox ', 'ybox ', 'width ', 'height', 'onpix ', 'xbar ',
      'ybar ', 'x2bar', 'y2bar ', 'xybar ', 'x2ybar', 'xy2bar', 'xedge ',
      'xedgey', 'yedge ', 'yedgey'],
      dtype='object')
```

```
In [4]: # let's 'reindex' the column names
letters.columns = ['letter', 'xbox', 'ybox', 'width', 'height', 'onpix', 'xbar',
                  'ybar', 'x2bar', 'y2bar', 'xybar', 'x2ybar', 'xy2bar', 'xedge',
                  'xedgey', 'yedge', 'yedgey']
print(letters.columns)

Index(['letter', 'xbox', 'ybox', 'width', 'height', 'onpix', 'xbar', 'ybar',
      'x2bar', 'y2bar', 'xybar', 'x2ybar', 'xy2bar', 'xedge', 'xedgey',
      'yedge', 'yedgey'],
      dtype='object')
```

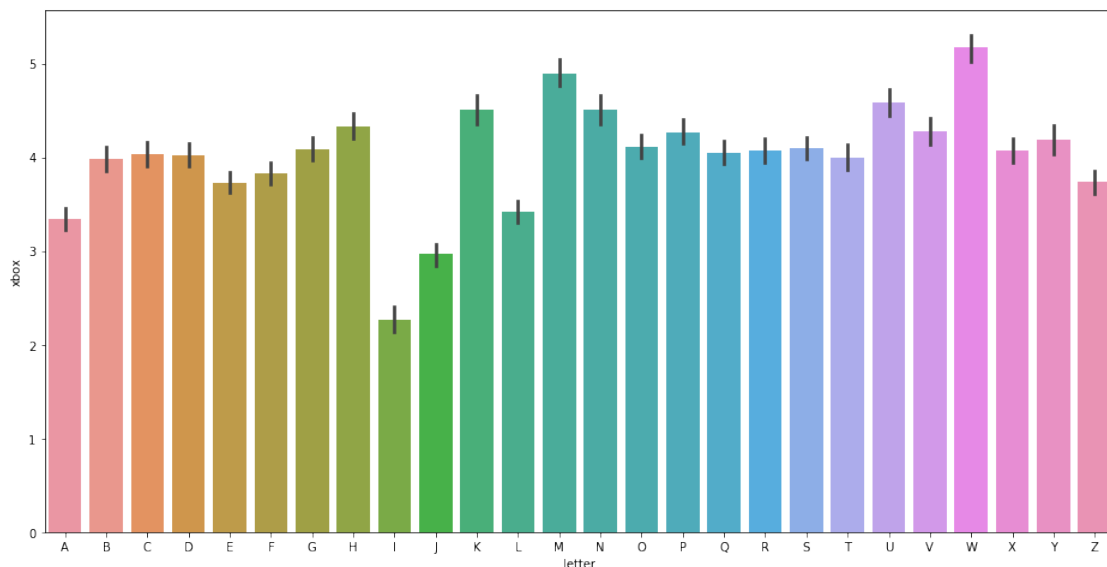
```
In [5]: order = list(np.sort(letters['letter'].unique()))
print(order)

['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z']
```

```
In [6]: # basic plots: How do various attributes vary with the letters
```

```
plt.figure(figsize=(16, 8))
sns.barplot(x='letter', y='xbox',
            data=letters,
            order=order)
```

```
Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x104cca5f8>
```



```
In [7]: letter_means = letters.groupby('letter').mean()
letter_means.head()
```

```

Out [7]:
      xbox    ybox    width    height    onpix    xbar    ybar  \
letter
A      3.337136  6.975919  5.128010  5.178707  2.991128  8.851711  3.631179
B      3.985640  6.962141  5.088773  5.169713  4.596606  7.671018  7.062663
C      4.031250  7.063859  4.701087  5.296196  2.775815  5.437500  7.627717
D      4.023602  7.244720  5.170186  5.288199  4.026087  7.539130  6.806211
E      3.727865  6.944010  4.756510  5.201823  3.679688  5.966146  7.352865

      x2bar    y2bar    xybar    x2ybar    xy2bar    xedge    xedgey  \
letter
A      2.755387  2.043093  7.802281  2.338403  8.465146  2.771863  6.321926
B      5.366841  5.571802  7.954308  5.506527  6.652742  3.117493  7.919060
C      5.927989  7.177989  8.773098  7.494565  11.947011  1.991848  8.876359
D      5.921739  6.508075  8.166460  5.111801  5.750311  3.365217  7.813665
E      4.223958  7.585938  8.507812  6.242188  10.341146  2.127604  8.298177

      yedge    yedgex
letter
A      2.875792  7.468948
B      6.612272  9.100522
C      4.080163  8.555707
D      3.971429  7.628571
E      6.022135  8.506510

```

1.2 Data Preparation

Let's conduct some data preparation steps before modeling. Firstly, let's see if it is important to **rescale** the features, since they may have varying ranges. For example, here are the average values:

```

In [9]: # average feature values
        round(letters.drop('letter', axis=1).mean(), 2)

```

```

Out [9]: xbox    4.02
        ybox     7.04
        width    5.12
        height   5.37
        onpix    3.51
        xbar     6.90
        ybar     7.50
        x2bar    4.63
        y2bar    5.18
        xybar    8.28
        x2ybar   6.45
        xy2bar   7.93
        xedge    3.05
        xedgey   8.34
        yedge    3.69

```

```
yedgex      7.80
dtype: float64
```

In this case, the average values do not vary a lot (e.g. having a diff of an order of magnitude). Nevertheless, it is better to rescale them.

```
In [10]: # splitting into X and y
X = letters.drop("letter", axis = 1)
y = letters['letter']

In [11]: # scaling the features
X_scaled = scale(X)

# train test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size = 0.3, ran
```

1.3 Model Building

Let's first build two basic models - linear and non-linear with default hyperparameters, and compare the accuracies.

```
In [12]: # linear model

model_linear = SVC(kernel='linear')
model_linear.fit(X_train, y_train)

# predict
y_pred = model_linear.predict(X_test)

In [13]: # confusion matrix and accuracy

# accuracy
print("accuracy:", metrics.accuracy_score(y_true=y_test, y_pred=y_pred), "\n")

# cm
print(metrics.confusion_matrix(y_true=y_test, y_pred=y_pred))
```

```
accuracy: 0.8523333333333334
```

```
[[198  0  0  0  0  0  1  1  0  1  1  1  0  0  0  0  0  1
   0  1  1  0  0  0  3  0]
 [  0 188  0  3  0  1  3  3  1  0  1  0  0  2  0  1  1  9
   3  0  0  1  0  1  0  0]
 [  1  0 200  0  7  0 12  1  0  0  5  0  0  0  3  0  0  0
   0  0  1  0  0  0  0  0]
 [  1 15  0 210  0  1  2  2  1  0  1  0  1  5  3  0  0  5
   0  1  0  0  0  0  0  0]
 [  0  1  3  0 204  2  6  1  0  0  1  5  0  0  0  0  2  2
   1  2  0  0  0  1  0  3]
```

```

[ 0  0  0  1  1 201  1  2  1  1  0  0  0  2  0  2  0  0
  3  7  1  0  1  0  1  0]
[ 0  1  9  4  2  2 167  1  0  1  4  3  1  0  1  0  9  1
  8  0  0  2  3  0  0  0]
[ 0  7  3 11  0  4  3 141  0  2  4  1  2  0 12  0  4 12
  0  0  4  2  0  4  1  0]
[ 0  0  2  3  0  6  0  0 184  9  0  0  0  0  1  0  0  0
  3  0  0  0  0  4  0  3]
[ 2  0  0  3  0  2  0  2 10 187  0  0  0  1  2  0  0  1
  5  0  1  0  0  0  0  4]
[ 0  1  5  2  0  0  1  3  0  0 198  2  2  0  0  0  0 19
  0  0  0  0  0 12  0  0]
[ 2  1  3  2  5  0  8  1  0  0  1 206  0  0  0  0  5  0
  2  1  0  0  0  0  0  0]
[ 0  3  0  0  0  0  0  3  0  0  0  0 222  1  0  0  0  2
  0  0  0  0  3  0  0  0]
[ 1  0  0  4  0  0  0  6  0  0  0  0  1 235  1  1  0  0
  0  0  0  1  1  0  0  0]
[ 3  0  4  7  0  0  0 21  0  0  0  0  2  0 163  3  2  3
  0  0  3  0 10  0  0  0]
[ 0  2  0  2  0 16  5  1  0  1  3  0  0  0  1 225  0  0
  0  0  0  1  0  0  8  0]
[ 3  1  0  0  4  0  9  0  0  2  0  1  0  0  6  0 198  0
  8  0  0  0  1  0  0  2]
[11 11  0  2  0  1  6  3  0  0 10  0  0  3  4  0  2 188
  0  1  0  0  0  1  0  0]
[ 1 13  0  0  9  5  8  0  7  1  0  2  0  0  0  1  6  1
155  6  0  0  0  3  0 10]
[ 0  0  0  1  1  4  2  4  1  0  0  0  0  0  0  0  0  1
  3 214  0  0  0  1  1  6]
[ 2  0  1  2  0  0  0  2  0  0  0  0  1  1  3  0  0  0
  0  0 211  0  1  0  0  0]
[ 2  2  0  0  0  0  1  3  0  0  0  0  0  1  0  1  0  3
  0  0  0 190  6  0  2  0]
[ 0  0  0  0  0  0  1  0  0  0  0  0  6  1  2  0  0  0
  0  0  0  0 212  0  0  0]
[ 0  2  0  4  5  0  1  0  2  3  3  3  0  0  1  0  0  1
  2  2  1  0  0 212  1  1]
[ 2  0  0  0  0  2  0  1  0  0  0  0  1  0  0  0  3  0
  0  4  1 10  0  2 211  0]
[ 1  0  0  0  3  0  0  0  0  6  0  0  0  0  0  0  5  0
18  1  0  0  0  1  0 194]]

```

The linear model gives approx. 85% accuracy. Let's look at a sufficiently non-linear model with randomly chosen hyperparameters.

```

In [14]: # non-linear model
          # using rbf kernel, C=1, default value of gamma

```

```

# model
non_linear_model = SVC(kernel='rbf')

# fit
non_linear_model.fit(X_train, y_train)

# predict
y_pred = non_linear_model.predict(X_test)

```

In [15]: # confusion matrix and accuracy

```

# accuracy
print("accuracy:", metrics.accuracy_score(y_true=y_test, y_pred=y_pred), "\n")

# cm
print(metrics.confusion_matrix(y_true=y_test, y_pred=y_pred))

```

accuracy: 0.9383333333333334

```

[[205  0  0  1  0  0  0  0  0  1  0  0  0  0  0  0  0
  0  0  0  0  0  0  2  0]
 [  0 205  0  3  1  0  1  0  0  0  0  0  0  0  0  0  0  6
  1  0  0  0  0  1  0  0]
 [  0  0 213  0  5  0  7  1  0  0  0  0  0  0  4  0  0  0
  0  0  0  0  0  0  0  0]
 [  0  4  0 234  0  0  1  3  0  0  0  0  0  3  1  0  0  2
  0  0  0  0  0  0  0  0]
 [  0  0  0  0 221  1  9  0  0  0  0  0  0  0  0  0  0  1
  0  0  0  0  0  0  0  2]
 [  0  0  0  1  0 215  1  1  1  0  0  0  0  1  0  1  0  0
  1  3  0  0  0  0  0  0]
 [  0  0  3  4  1  1 202  0  0  0  0  1  1  0  2  0  0  1
  0  0  0  1  2  0  0  0]
 [  0  7  0  5  0  0  4 177  0  0  2  0  1  0  3  0  4 13
  0  0  1  0  0  0  0  0]
 [  0  0  1  1  0  3  0  0 194 11  0  0  0  0  0  1  0  0
  2  0  0  0  0  2  0  0]
 [  1  0  0  1  0  0  0  0  6 206  0  0  0  1  2  0  0  0
  2  0  0  0  0  1  0  0]
 [  0  4  0  2  0  0  0  4  0  0 217  0  1  0  0  0  0 14
  0  0  0  0  0  3  0  0]
 [  0  0  1  0  2  0  6  0  0  0  1 222  0  0  0  0  0  3
  0  0  0  0  0  2  0  0]
 [  0  5  0  0  0  0  0  2  0  0  0  0 225  0  0  0  0  0
  0  0  0  0  2  0  0  0]
 [  0  2  0  1  0  0  0  2  0  0  0  0  1 239  3  0  0  2
  0  0  0  0  0  0  1  0]

```

```
[ 0  0  0  1  0  0  0  0  0  0  0  0  0  0 209  0  1  1
  0  0  1  0  8  0  0  0]
[ 0  2  0  3  3 11  1  1  0  0  0  0  0  0  1 237  1  0
  0  0  0  0  0  0  5  0]
[ 0  0  0  0  2  0  2  0  0  0  0  0  0  0  6  0 222  0
  1  0  0  0  2  0  0  0]
[ 0 10  0  2  0  0  0  0  0  0  1  0  0  4  0  0  2 224
  0  0  0  0  0  0  0  0]
[ 0  3  0  0  2  3  0  0  0  0  0  0  0  0  0  0  0  0
 220  0  0  0  0  0  0  0]
[ 0  0  0  1  0  2  0  2  0  0  0  0  0  0  0  1  0  1
  0 228  0  0  0  3  1  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  1  1  0  0  0
  0  0 222  0  0  0  0  0]
[ 0  7  0  0  0  0  0  1  0  0  0  0  1  4  0  1  0  0
  0  0  0 193  1  0  3  0]
[ 0  1  0  0  0  0  1  1  0  0  0  0  0  1  0  0  0  0
  0  0  1  0 217  0  0  0]
[ 0  2  0  3  2  0  0  0  1  0  2  0  0  0  0  0  0  1
  0  0  0  0  0 233  0  0]
[ 2  0  0  0  0  0  0  0  0  0  0  0  1  0  0  0  1  0
  0  1  2  2  0  0 228  0]
[ 0  0  0  0  2  0  0  0  0  0  0  0  0  0  0  0  4  0
  1  0  0  0  0  0  0 222]]
```

The non-linear model gives approx. 93% accuracy. Thus, going forward, let's choose hyperparameters corresponding to non-linear models.

1.4 Grid Search: Hyperparameter Tuning

Let's now tune the model to find the optimal values of C and gamma corresponding to an RBF kernel. We'll use 5-fold cross validation.

```
In [16]: # creating a KFold object with 5 splits
         folds = KFold(n_splits = 5, shuffle = True, random_state = 101)

         # specify range of hyperparameters
         # Set the parameters by cross-validation
         hyper_params = [ {'gamma': [1e-2, 1e-3, 1e-4],
                           'C': [1, 10, 100, 1000]}]

         # specify model
         model = SVC(kernel="rbf")

         # set up GridSearchCV()
         model_cv = GridSearchCV(estimator = model,
```



```

        param_grid = hyper_params,
        scoring= 'accuracy',
        cv = folds,
        verbose = 1,
        return_train_score=True)

    # fit the model
    model_cv.fit(X_train, y_train)

```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

[Parallel(n_jobs=1)]: Done 60 out of 60 | elapsed: 10.0min finished

```

Out[16]: GridSearchCV(cv=KFold(n_splits=5, random_state=101, shuffle=True),
    error_score='raise',
    estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False),
    fit_params=None, iid=True, n_jobs=1,
    param_grid=[{'gamma': [0.01, 0.001, 0.0001], 'C': [1, 10, 100, 1000]}],
    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
    scoring='accuracy', verbose=1)

```

```

In [17]: # cv results
cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results

```

```

Out[17]:
   mean_fit_time  mean_score_time  mean_test_score  mean_train_score  param_C \
0         2.531680         1.421701         0.832714         0.844679         1
1         5.496493         1.896336         0.677214         0.683821         1
2        10.647400         2.109896         0.217571         0.227982         1
3         1.609691         0.936208         0.911214         0.929304        10
4         2.501903         1.427731         0.808929         0.817857        10
5         5.405686         1.914143         0.677000         0.683643        10
6         2.070304         0.750065         0.947786         0.978411       100
7         1.858425         0.965986         0.863357         0.878768       100
8         2.490351         1.424067         0.805714         0.813179       100
9         2.196414         0.660221         0.951714         0.997357      1000
10        2.385556         0.784292         0.898357         0.929982      1000
11        2.175660         1.081910         0.844357         0.860679      1000

   param_gamma  params  rank_test_score \
0         0.01  {'C': 1, 'gamma': 0.01}      7
1         0.001 {'C': 1, 'gamma': 0.001}     10
2         0.0001 {'C': 1, 'gamma': 0.0001}     12
3         0.01  {'C': 10, 'gamma': 0.01}       3

```

4	0.001	{'C': 10, 'gamma': 0.001}	8
5	0.0001	{'C': 10, 'gamma': 0.0001}	11
6	0.01	{'C': 100, 'gamma': 0.01}	2
7	0.001	{'C': 100, 'gamma': 0.001}	5
8	0.0001	{'C': 100, 'gamma': 0.0001}	9
9	0.01	{'C': 1000, 'gamma': 0.01}	1
10	0.001	{'C': 1000, 'gamma': 0.001}	4
11	0.0001	{'C': 1000, 'gamma': 0.0001}	6

	split0_test_score	split0_train_score	...	split2_test_score \
0	0.826429	0.843929	...	0.828571
1	0.684643	0.682054	...	0.682500
2	0.203929	0.212679	...	0.206429
3	0.914643	0.929643	...	0.902500
4	0.805714	0.818125	...	0.803571
5	0.684286	0.681875	...	0.682500
6	0.948571	0.978482	...	0.941429
7	0.861071	0.879554	...	0.850357
8	0.801071	0.813571	...	0.803929
9	0.953571	0.997054	...	0.948571
10	0.897500	0.931339	...	0.889643
11	0.844643	0.860982	...	0.829286

	split2_train_score	split3_test_score	split3_train_score \
0	0.847679	0.843929	0.843571
1	0.687946	0.670000	0.679554
2	0.228393	0.237500	0.240625
3	0.930357	0.919286	0.928929
4	0.819643	0.820714	0.814375
5	0.687500	0.670000	0.680089
6	0.978393	0.952857	0.978661
7	0.882143	0.875357	0.875536
8	0.814286	0.817143	0.810000
9	0.997411	0.952857	0.997321
10	0.931607	0.908571	0.927589
11	0.864018	0.856071	0.858304

	split4_test_score	split4_train_score	std_fit_time	std_score_time \
0	0.830000	0.844018	0.036476	0.047152
1	0.676429	0.686429	0.186916	0.054003
2	0.232857	0.240625	0.260968	0.097644
3	0.911071	0.928125	0.022005	0.020297
4	0.805000	0.819464	0.024353	0.010922
5	0.677143	0.685446	0.216667	0.111240
6	0.949286	0.978571	0.088442	0.041193
7	0.863929	0.879464	0.047128	0.009356
8	0.800000	0.815357	0.014262	0.063285
9	0.953929	0.997232	0.022609	0.002775

10	0.900357	0.931696	0.040057	0.026058
11	0.843571	0.861607	0.019762	0.006161

	std_test_score	std_train_score
0	0.006221	0.001514
1	0.005622	0.003021
2	0.014490	0.011507
3	0.005654	0.000746
4	0.006227	0.001897
5	0.005788	0.002608
6	0.003742	0.000249
7	0.008074	0.002262
8	0.006140	0.001815
9	0.002183	0.000237
10	0.006197	0.001921
11	0.008717	0.002125

[12 rows x 22 columns]

```
In [18]: # converting C to numeric type for plotting on x-axis
cv_results['param_C'] = cv_results['param_C'].astype('int')

# # plotting
plt.figure(figsize=(16,6))

# subplot 1/3
plt.subplot(131)
gamma_01 = cv_results[cv_results['param_gamma']==0.01]

plt.plot(gamma_01["param_C"], gamma_01["mean_test_score"])
plt.plot(gamma_01["param_C"], gamma_01["mean_train_score"])
plt.xlabel('C')
plt.ylabel('Accuracy')
plt.title("Gamma=0.01")
plt.ylim([0.60, 1])
plt.legend(['test accuracy', 'train accuracy'], loc='upper left')
plt.xscale('log')

# subplot 2/3
plt.subplot(132)
gamma_001 = cv_results[cv_results['param_gamma']==0.001]

plt.plot(gamma_001["param_C"], gamma_001["mean_test_score"])
plt.plot(gamma_001["param_C"], gamma_001["mean_train_score"])
plt.xlabel('C')
plt.ylabel('Accuracy')
plt.title("Gamma=0.001")
plt.ylim([0.60, 1])
```

```
plt.legend(['test accuracy', 'train accuracy'], loc='upper left')
plt.xscale('log')
```

```
# subplot 3/3
```

```
plt.subplot(133)
```

```
gamma_0001 = cv_results[cv_results['param_gamma']==0.0001]
```

```
plt.plot(gamma_0001["param_C"], gamma_0001["mean_test_score"])
```

```
plt.plot(gamma_0001["param_C"], gamma_0001["mean_train_score"])
```

```
plt.xlabel('C')
```

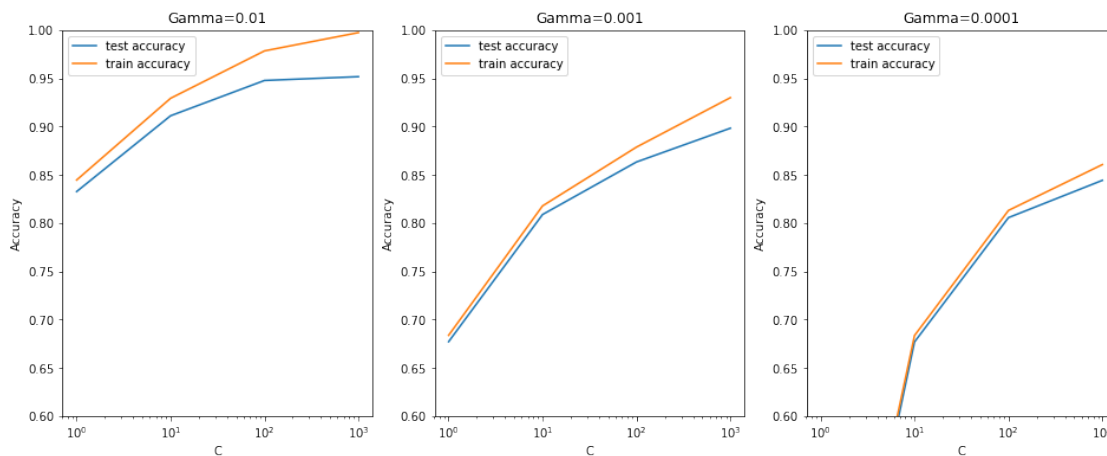
```
plt.ylabel('Accuracy')
```

```
plt.title("Gamma=0.0001")
```

```
plt.ylim([0.60, 1])
```

```
plt.legend(['test accuracy', 'train accuracy'], loc='upper left')
```

```
plt.xscale('log')
```



The plots above show some useful insights: - Non-linear models (high gamma) perform *much better* than the linear ones - At any value of gamma, a high value of C leads to better performance - None of the models tend to overfit (even the complex ones), since the training and test accuracies closely follow each other

This suggests that the problem and the data is **inherently non-linear** in nature, and a complex model will outperform simple, linear models in this case.

Let's now choose the best hyperparameters.

```
In [19]: # printing the optimal accuracy score and hyperparameters
```

```
best_score = model_cv.best_score_
```

```
best_hyperparams = model_cv.best_params_
```

```
print("The best test score is {0} corresponding to hyperparameters {1}".format(best_s
```

The best test score is 0.9517142857142857 corresponding to hyperparameters {'C': 1000, 'gamma'

1.4.1 Building and Evaluating the Final Model

Let's now build and evaluate the final model, i.e. the model with highest test accuracy.

In [20]: *# model with optimal hyperparameters*

```
# model
model = SVC(C=1000, gamma=0.01, kernel="rbf")

model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# metrics
print("accuracy", metrics.accuracy_score(y_test, y_pred), "\n")
print(metrics.confusion_matrix(y_test, y_pred), "\n")
```

accuracy 0.9596666666666667

```
[[206  0  0  0  0  0  0  0  0  0  0  1  0  0  0  0  0
   0  0  0  0  0  0  2  0]
 [  0 211  0  2  1  0  1  0  0  0  0  0  0  0  0  0  0  1
   2  0  0  0  0  0  0  0]
 [  0  0 220  0  3  0  4  1  0  0  0  0  0  0  2  0  0  0
   0  0  0  0  0  0  0  0]
 [  0  3  0 236  0  1  0  1  0  1  0  0  0  2  2  0  0  1
   1  0  0  0  0  0  0  0]
 [  0  0  1  0 225  1  4  0  0  0  0  0  0  0  0  0  0  0
   1  0  0  0  0  0  0  2]
 [  0  0  0  1  0 217  0  0  1  1  0  0  0  1  0  3  0  0
   0  0  0  1  0  0  0  0]
 [  0  0  2  3  1  0 209  0  0  0  0  0  1  0  1  0  0  0
   0  0  0  1  1  0  0  0]
 [  0  1  3  5  0  0  2 195  1  1  2  1  1  0  1  0  3  1
   0  0  0  0  0  0  0  0]
 [  0  0  0  1  0  1  0  0 203  8  0  0  0  0  0  0  1  0
   0  0  0  0  0  1  0  0]
 [  0  0  0  0  0  0  0  0  7 209  0  0  0  1  0  0  0  1
   0  0  2  0  0  0  0  0]
 [  0  1  0  0  2  0  0  5  0  0 228  0  0  0  0  0  0  5
   0  0  0  0  0  4  0  0]
 [  0  0  0  0  0  0  1  1  0  0  0 232  0  0  0  1  0  1
   0  1  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0  0  0  0  0 230  2  0  0  0  0
   0  0  0  0  2  0  0  0]
 [  0  3  0  1  0  0  0  0  0  0  0  0  1 244  0  0  0  1
   0  0  0  0  1  0  0  0]
 [  0  0  2  0  0  0  2  0  0  1  0  0  2  0 210  0  1  1
   0  0  1  0  1  0  0  0]
 [  0  0  0  0  1  8  0  1  1  0  0  0  0  0  0 252  1  0
```

```

    0  0  0  0  0  0  1  0]
[ 0  0  0  0  3  0  1  0  0  0  0  0  0  0  2  1 226  0
  0  0  0  0  0  0  0  2]
[ 0  8  0  1  0  0  1  3  0  0  3  0  0  3  0  0  0 224
  0  0  0  0  0  0  0  0]
[ 0  1  0  0  1  2  1  0  0  0  0  0  0  0  0  0  0  0
223  0  0  0  0  0  0  0]
[ 0  0  0  1  0  0  0  1  0  0  0  0  0  0  0  1  0  0
0 235  0  0  0  1  0  0]
[ 2  0  0  0  0  0  1  2  0  0  0  0  0  2  0  0  0  0
  0  0 217  0  0  0  0  0]
[ 0  4  0  0  0  0  0  0  0  0  0  0  1  0  0  1  0  0
  0  0  0 203  1  0  1  0]
[ 0  0  0  0  0  0  1  0  0  0  0  0  0  0  0  1  0  0
  0  0  0  0 220  0  0  0]
[ 0  1  0  3  2  0  0  0  0  0  3  1  0  0  0  0  0  1
  0  0  0  0  0 232  1  0]
[ 1  0  0  0  0  0  0  0  0  1  0  0  0  0  0  0  0  0
  1  0  1  2  1  0 230  0]
[ 0  0  0  0  2  1  0  0  0  0  0  0  0  0  0  0  3  0
  1  1  0  0  0  0  0 221]]

```

1.5 Conclusion

The accuracy achieved using a non-linear kernel (~ 0.95) is much higher than that of a linear one (~ 0.85). We can conclude that the problem is highly non-linear in nature.